

Assignment Project Exam Help
Synchronization

<https://tutorcs.com>

WeChat: cstutorcs
Lecturer: Dr. Joseph Doyle

Introduction

- Events occur at different nodes in the network
- These events are recorded using different clocks
- How do other nodes in the network determine when events took place

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Why is it important

- May seem like a trivial issue but actually very important
- How can we determine if data is up-to-date unless we can determine the order of writes to the data
- Is it possible to sell stock that does not exist?

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Synchronize Clocks

- Simplest solution would be to synchronise clocks on every node in the network
- In practice this is very difficult to achieve

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Hardware Clocks

- Hardware clocks in computers are crystal oscillators which are connected to oscillation counter circuitry
- This counter is then scaled to provide an approximate representation of physical time

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estutorcs

Problems with Hardware Clocks

- Clocks are not all exactly the same
- The clock rate is affected by
 - Physical differences in the crystal
 - Temperature differences
 - Voltage
 - Humidity

Assignment Project Exam Help

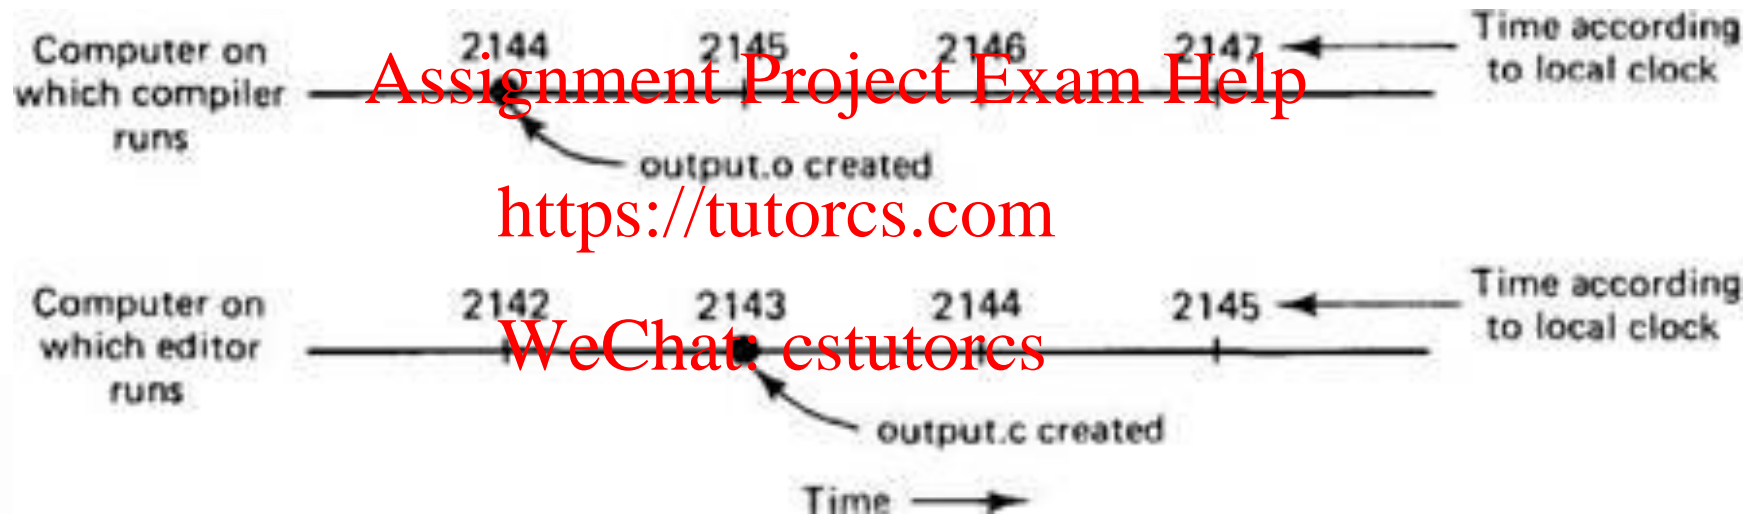
<https://tutorcs.com>

WeChat: cstutorcs

Problems with Hardware Clocks

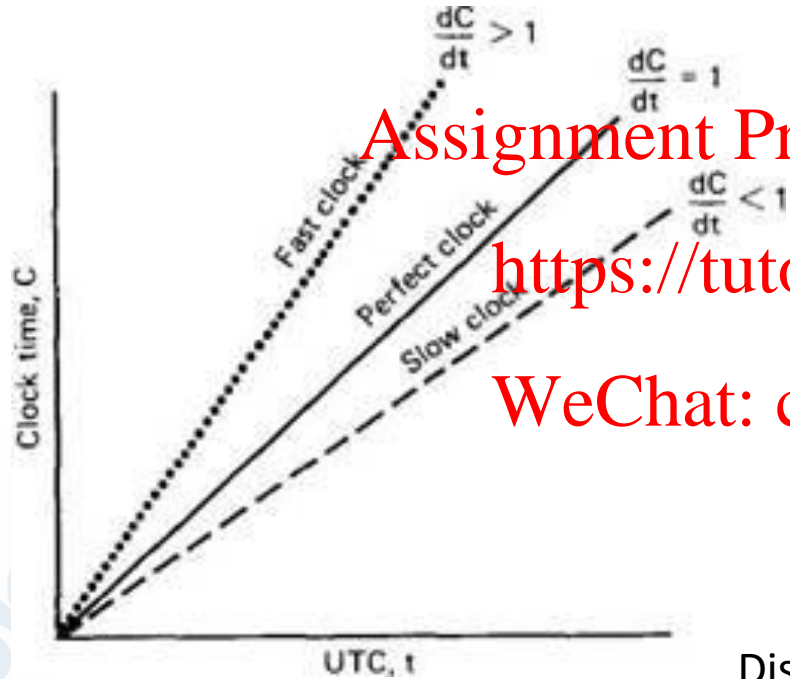
- This results in three problems:
 - **Skew:** This is a disagreement in the reading of two clocks
 - **Drift:** This is the difference in the rate at which two clocks count the time
 - **Clock Drift Rate:** This is the difference in precision between a hardware clock and a perfect reference clock. Approximately 10^{-6} sec/sec in normal clocks and 10^{-7} - 10^{-8} sec/sec in high precision clocks

Clock Synchronization



Distributed Operating Systems. Andrew Tanenbaum

Clock Synchronization



- Assuming a normal clock drift rate of 10^{-6} sec/sec a drift of 1ms occurs in approximately 17 minutes and a drift of 1s occurs in approximately 12 days

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Distributed Operating Systems. Andrew Tanenbaum

Synchronizing Clocks

- External Synchronization

- An external authoritative clock updates the clock of nodes in the network
- Skew is limited to the delay between the authoritative clock server and the node

- Internal Synchronization

- Nodes in the network collaborate to limit the skew to a delay bound
- Delay is larger here as a round trip is required

Software Based Clock Synchronization

- Christian's Algorithm
 - Berkley Algorithm
 - Network Time Protocol (Internet)
- Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

Christian's Algorithm

- Based around the observation that round trip time in LAN networks is sufficiently small to ensure reasonable clock accuracy
- It requires a clock server which is synchronized to UTC time

Christian's Algorithm

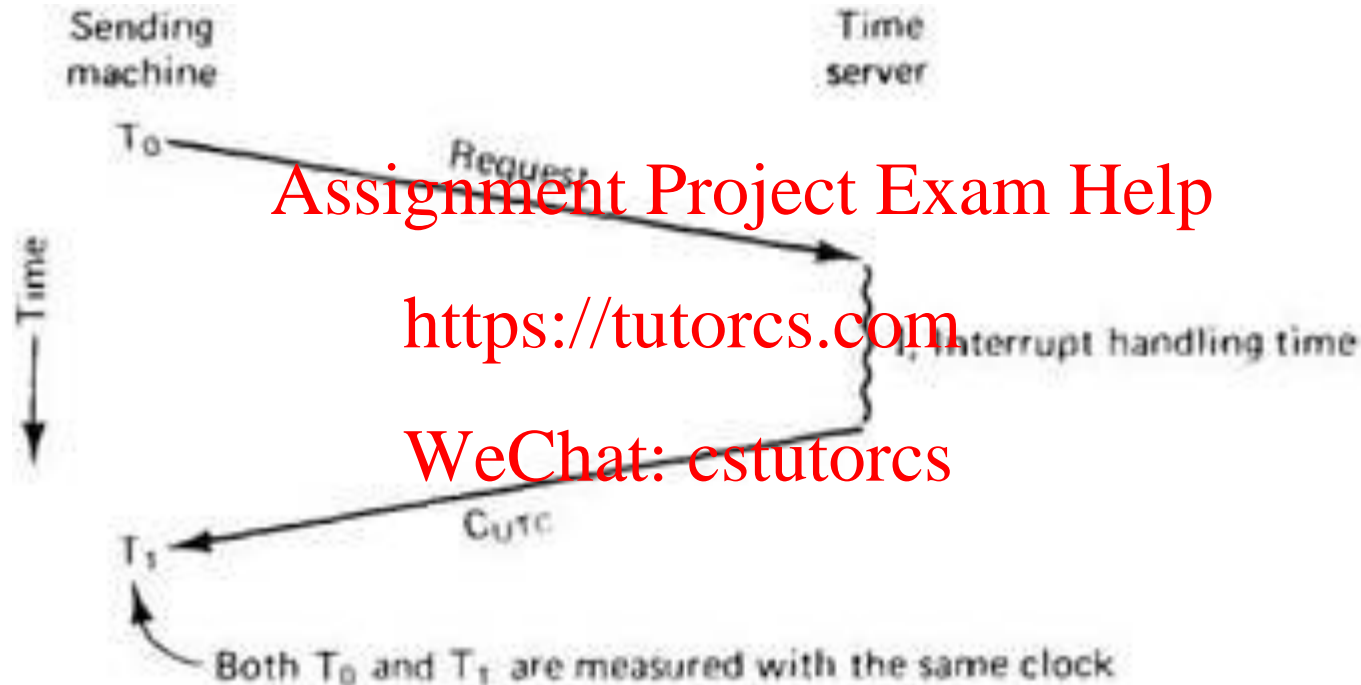
- A node in the network sends a request to the clock server and records the round trip time RTT
- The clock server measures its current time T and sends this to the node
- The node then updates its time t to be $t = T + RTT/2$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Christian's Algorithm



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Distributed Operating Systems. Andrew Tanenbaum

Christian's Algorithm Problems

- Assumes that the duration of the two parts of the round trip are equal
- Not really suitable outside a LAN as the RTT can increase dramatically
- The clock server is a central point of failure

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Berkeley Algorithm

- Variation of Christians Algorithm which does not require a clock server that is synchronized to UTC time <https://tutorcs.com>
- Uses a manager server instead of a clock server to alter the clocks of nodes in the network

Berkeley Algorithm


- Manager server periodically polls all nodes in the network
- It records the RTT and uses this to estimate the clock of the node in a similar fashion to Christian's algorithm
- It averages the values obtained from all nodes
- It instructs the nodes to alter their clocks based upon this average to synchronize the times of the nodes
- If the manager fails a new manager is elected using a manager election algorithm (This is discussed later)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estutotors

Berkeley Algorithm



Distributed Operating Systems. Andrew Tanenbaum

Berkeley Algorithm Problems

- Again this is really only suitable for LAN networks
- Also while the nodes are synchronized with each other they are not synchronized with external systems if this is required as is the case with Christian's algorithm

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Network Time Protocol (NTP)

- The goal of this protocol is to improve on the previously discussed algorithms to allow:
 - The ability to synchronize clients via the Internet
 - Reliable service in the event of lengthy losses of connectivity
 - Provide protection against interference

Assignment Project Exam Help

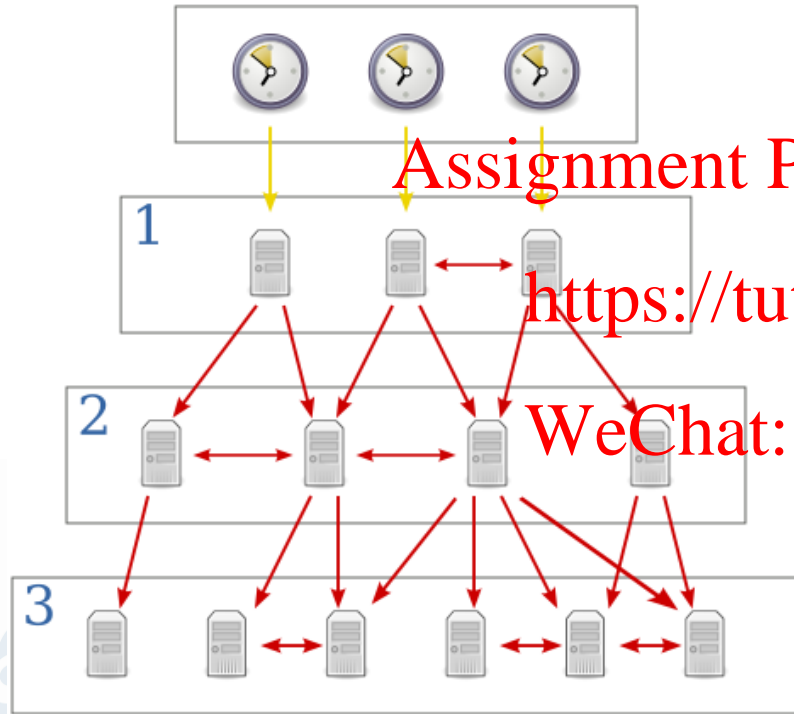
<https://tutorcs.com>

WeChat: estutorcs

Network Time Protocol (NTP)

- NTP uses a hierarchical model with different stratum in the hierarchy assigned different levels of accuracy
 - **Stratum 0** are high precision timekeeping devices such as atomic clocks, GPS or other radio clocks. They are known as reference clocks
 - **Stratum 1** are computers which are synchronised to within a few microseconds of a Stratum 0 devices. These are the primary time servers.
 - **Stratum 2** are nodes connected to primary time servers. It is possible to connect to multiple primary time servers

Network Time Protocol (NTP)



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

https://en.wikipedia.org/wiki/Network_Time_Protocol

Network Time Protocol (NTP)

- Up to 15 stratum are possible in NTP
- The higher the stratum number the less accurate the clock of the node due to the increased RTT

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

NTP Timestamps

- Uses a 64 bit timestamp with 32 bits for seconds and 32 bits for fractional seconds
- It has a theoretical resolution of 233 picoseconds
- It has a timescale that rolls over every 136 years. First rollover will occur on February 7th 2036 (Something to look forward to)



/QMUL



@QMUL

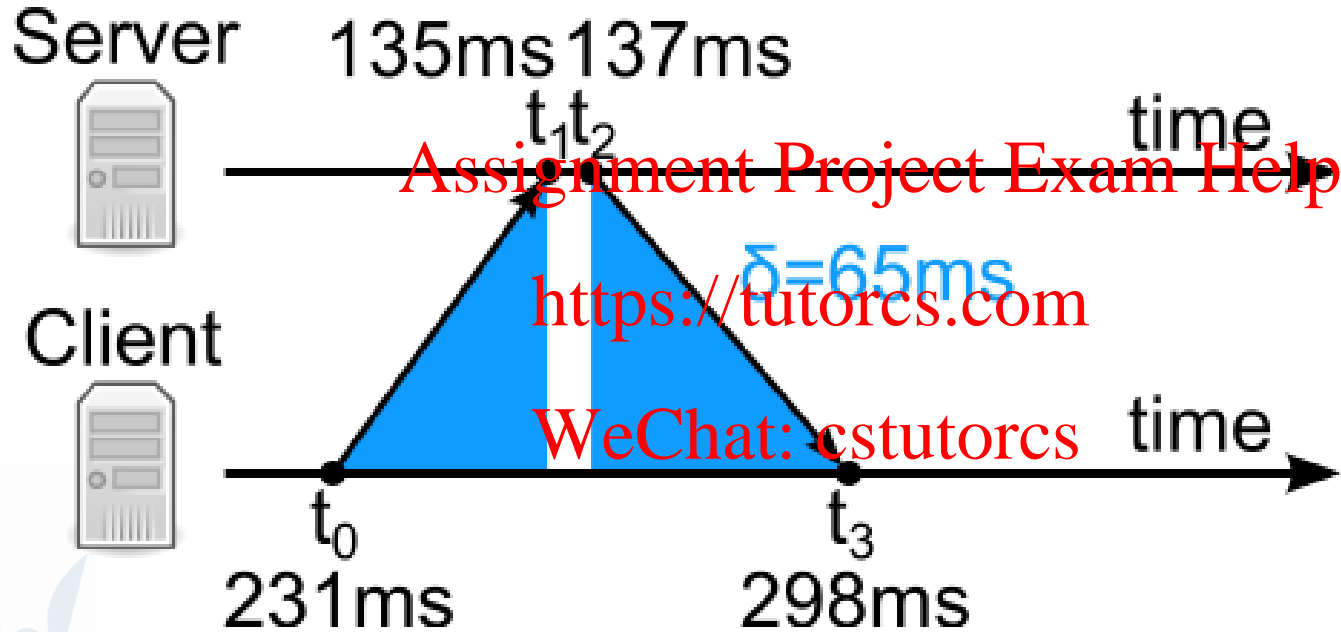
NTP Algorithm

- The NTP client will regularly poll one or more server
- Using this information they will calculate the time offset θ and round trip delay δ as:

$$- \theta = \frac{(t_1 - t_0) + (t_2 - t_3)}{2}$$

$$- \delta = (t_3 - t_0) - (t_2 - t_1)$$

NTP Algorithm



https://en.wikipedia.org/wiki/Network_Time_Protocol

NTP Algorithm

- Once these values are returned to the client they are passed through a statistical filter to eliminate outliers
- An estimate of the offset is calculated from the best three remaining candidates (It favours messages from higher stratum)
- The client's clock is then adjusted gradually to reduce the offset

Assignment Project Exam Help


<https://tutorcs.com>

WeChat: cstutorcs

Synchronization Requirements

- **Causality:** real-time order \sim timestamp order
- **Groups/replicated:** All members of the group see events in the same order
<https://tutorcs.com>
- **Multiple-copy-updates:** order of the updates, consistency conflicts
WeChat: cstutorcs
- **Serializability of transactions:** Common order of transaction order

Happened-Before Relations ($a \rightarrow b$)



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- a and b are defined as events in the same process
- If a occurs before b then $a \rightarrow b$
- For example if a is a message being sent and b is a message being received then $a \rightarrow b$
- $a \parallel b$ if neither $a \rightarrow b$ and $b \rightarrow a$ (a and b are concurrent)
- If $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$


Lamport Timestamps

- The rules for this algorithm are as follows:
 - A process increments its counter before each event that it processes
 - The process includes this counter when it sends a message
 - On receiving a message the counter of the recipient is updated if necessary to the greater of the its current counter and timestamp received in its message. The counter is then incremented by one to indicate that the message has been received.

Lamport Timestamps


- If we define a process as p_i , an event as e , a counter as L_i and a timestamp as $L_i(e)$ we can define this algorithm as
- At p_i : before each event $L_i = L_i + 1$
- When p_i sends a message m to p_j
 - p_i : $L_i = L_i + 1$; $t = L_i$; $message = (m, t)$
 - p_j : $L_j = \max(L_j, t)$; $L_j = L_j + 1$
 - $L_j(receive\ event) = L_j$

Lamport Timestamps



(a)

(Uncorrected)



(b)

(Lamport Timestamps)

Distributed Operating Systems. Andrew Tanenbaum

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: estutorcs

Lamport Timestamps Problems

- Suppose there are two event a and b
- If there is any way that event a could have influenced event b then we can state that $a \rightarrow b$ and the Lamport timestamp of event a is less than the Lamport timestamp of event b
- However, if we have two Lamport timestamps $L(a) < L(b)$ we cannot explicitly state that $a \rightarrow b$ as Lamport timestamps do not fulfil the **strong clock consistency condition**
- Lamport timestamps only create a **partial ordering** of events
- They can be used to create a **total ordering** of events by including an arbitrary mechanism to break ties (albeit with the caveat that this cannot be used to imply a causal relationship)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Total Ordered Lamport Timestamps

- Expand Timestamp to $(L_i(e), i)$
- $(L_i(e), i) < (L_j(e), j)$
- If $L_i(e) < L_j(e)$ or $(L_i(e) = L_j(e) \text{ and } i < j)$

Assignment Project Exam Help

<https://tutorcs.com>


WeChat: cstutorcs

Total Order Multicasting Examples



Database is updated leaving an inconsistent state

Total Order Multicasting Examples



All receivers see all messages in the same order
(which is not necessarily the original sending order)
E.g. Group Updates

Vector Clocks

- Vector clocks expand on Lamport Timestamps by including additional data for detecting causality violations
- Each Process P_i maintains a vector V_i
- $V_i[i]$ is the number of events that have occurred at P_i
- If $V_i[j] = k$ is then P_i knows about the k events that have occurred at P_j

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Vector Clocks


- Order of timestamps
- $V = V'$ iff $V[j] = V'[j]$ for all j
- $V \leq V'$ iff $V[j] \leq V'[j]$ for all j
- $V < V'$ iff $V \leq V'$ and $V \neq V'$
- Order of events
- $e \rightarrow e' \Rightarrow V(e) < V(e')$
- $V(e) < V(e') \Rightarrow e \rightarrow e'$
- $e \parallel e'$ if not $V(e) \leq V(e')$
- and not $V(e') \leq V(e)$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Vector Clocks



https://en.wikipedia.org/wiki/Vector_clock

Vector Clocks

- The algorithm for vector clocks is as follows:
- P_i multicast $V_i[i] = V_i[i] + 1$
- Each message includes V_i
- For each P_j which is receiving a message: The message can be delivered when
 - $V_i[i] = V_j[i] + 1$ (All previous message from i have arrived)
 - $V_j[k] \geq V_i[k]$ for all $k, k \neq i$ (j has seen all the message i has seen when the message was sent)
- Upon delivery of message $V_j[i] = V_j[i] + 1$

Global State

- Timestamps can be used for a variety of purposes such saving a snapshot of a distributed system
- To create a snapshot the system needs information on:
 - The state of processes
 - Messages in transfer
- There are potential problems which can affect this snapshot namely:
 - Garbage Collection
 - Deadlock
 - Termination

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutores

Garbage Collection

- Garbage collection is an automatic form of memory management
- Very important on distributed systems as manual memory management is more difficult when using multiple systems
- For example consider a simple distributed messaging system
- Messages must be kept until they can be processed after which they can be discarded by a garbage collector
- Need to consider timestamps of messages to determine if they should be included in a snapshot or if they should be discarded

Assignment Project Exam Help

<https://tutorcs.com>

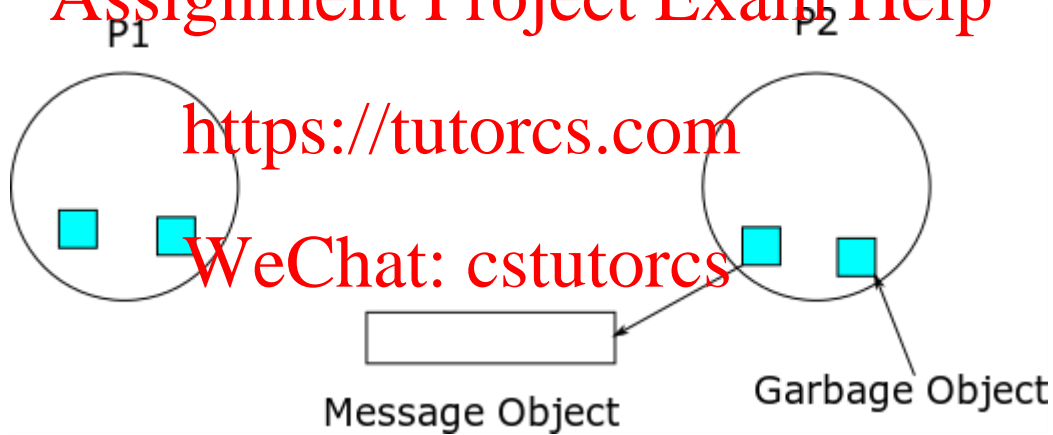
WeChat: ostutorcs

Garbage Collection

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Deadlock


- Deadlock occurs when two processes are waiting for each other to take an action such as releasing a lock which is associated with a resource
- This can occur for a variety of reasons. Consider the following example
- P_1 has a lock on resource R_1 and P_2 has a lock on resource R_2 . P_1 needs to obtain a lock on R_2 before it will release R_1 and P_2 needs to obtain a lock on R_1 before it will release R_2 .
- Unless action is taken to resolve the deadlock P_1 and P_2 will continue to wait for a lock they cannot obtain.
- It would be useful to resolve deadlocks before taking a snapshot as they prevent progress in distributed systems

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Deadlock



<https://en.wikipedia.org/wiki/Deadlock>

Global Snapshot

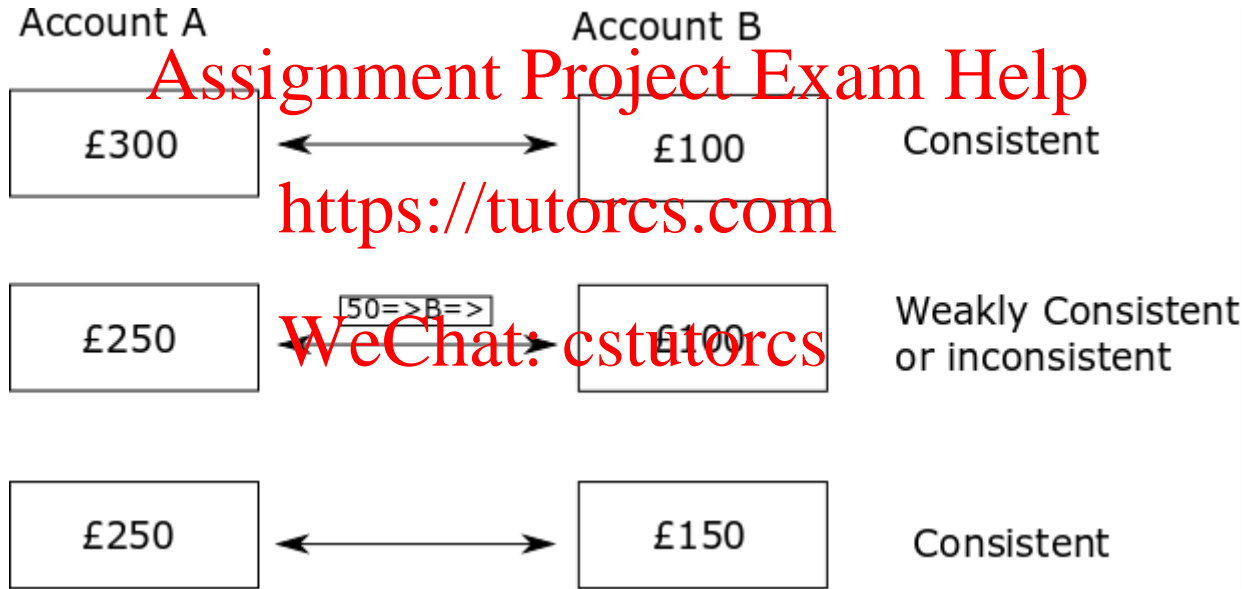
- To create a global snapshot the following algorithm can be used:
- At each node i
 - A local clock records the time T_i
 - A state S_i is recorded as a list of tuples containing {event, timestamp}
- The system state S is defined as a vector of the state of each node S_i
- A snapshot will contain information on each node up to time T
- A snapshot can be considered consistent or inconsistent

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Inconsistent Snapshot Example



Consistent Snapshot


- A snapshot is consistent if it contains all the event which happened before the snapshot time T
- Consider the example where changes to a local value x_i at processor i only need to be sent to another processor when they exceed a certain threshold (>19)
- A vector clock is incremented when x_i changes
- S records all changes to the x_i value

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Consistent Snapshot



Chandy Lamport Algorithm

- To create a consistent global state of a distributed system the Chandy-Lamport Algorithm can be used
- The algorithm assumes that:
 - There are no network failures and all messages arrive intact
 - The snapshot algorithm does not interfere with the normal operation of the processes
- It would be possible to modify this algorithm with TCP/IP to relax the no network failures assumption

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Chandy Lamport Algorithm

- The process which is initiating the snapshot process:
 - Saves its own state
 - Sends a snapshot request to all processes bearing a snapshot token
- When a process receives a snapshot request it
 - Sends the snapshot process its saved state
 - Attaches the snapshot token to all subsequent messages
- When a process receives a message that does not have the snapshot token when it has received the snapshot token it forwards the message to the snapshotting process so that it can be included in the snapshot

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Resource Reservation

- Timestamps are also useful for reserving resources
- The process which controls resource allocation can be centralized or distributed
- Centralized in general is easy to implement but are not fault tolerant
- Distributed processes more complicated but are in general more fault tolerant
- The other disadvantage of distributed solution is that it tends to be more difficult to debug
- This is sometimes referred to as Mutual Exclusion

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Centralized Resource Allocation

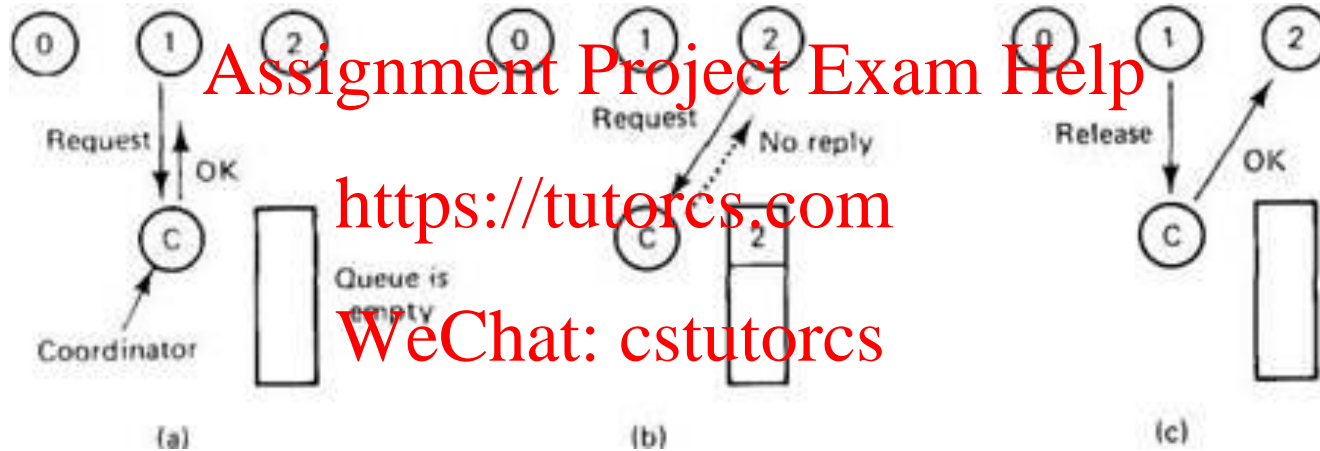
- In the central case a process will ask the coordinator for access to the resource
- The coordinator maintains a queue and if the queue is empty it will grant access to the process
- If another process contacts the coordinator it is placed in the queue and the coordinator does not respond to the request
- When the first process finishes it informs the coordinator and the coordinator then allows the next process in the queue access to the resource

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutores

Centralized Resource Allocation



Distributed Operating Systems. Andrew Tanenbaum

Centralized Resource Allocation

- The general requirements of mutual exclusion are:
- **Safety:** Only one process is allowed access to the resource at any given time (The algorithm fulfils this requirement)
- **Liveness:** All requests eventually succeed. Deadlock is prevented. (This requirement is not specifically fulfilled. Several methods can be used to break deadlock with the simplest being timeout functionality)
- **Fairness:** If request A happens before request B then A is honoured before B (Timestamps play a role in this requirement at request A may have been delayed in the network but it should still be honoured before request B)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutores

Distributed Resource Allocation

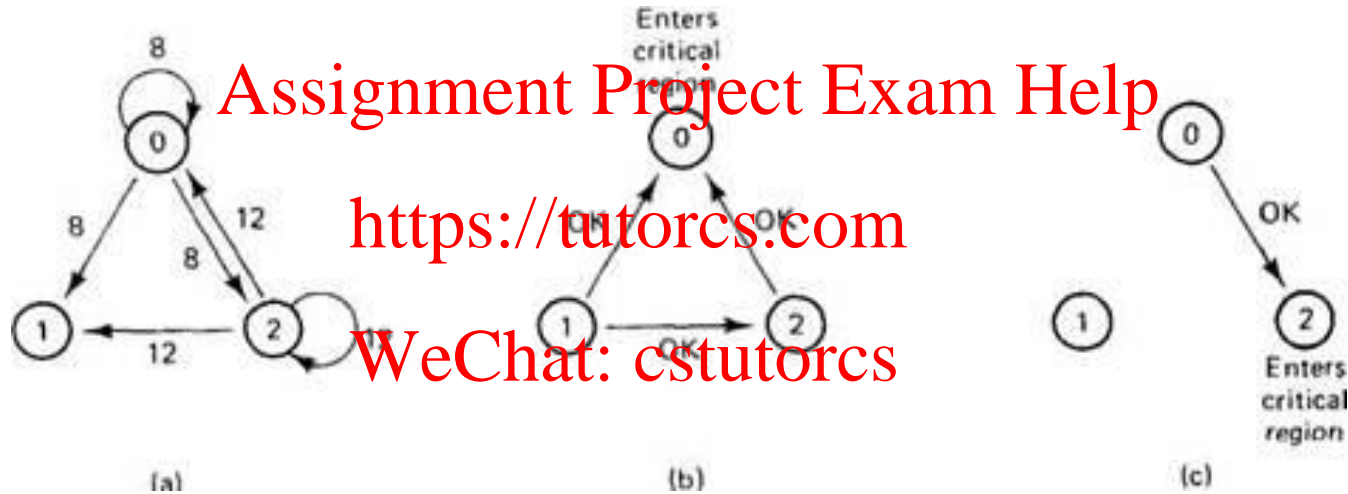
- Also possible to use a distributed scheme
- When a process wants a resource it multicasts a request to all process and waits for the response
- When a process receives a request if it does not want the response it responds immediately
- If it is using the resource or wants to use the resource and the timestamp of its request is lower than the received timestamp it puts the request in a queue
- When it finishes using the request it replies to all requests in the queue

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Distributed Resource Allocation



Distributed Operating Systems. Andrew Tanenbaum

Token Ring Resource Allocation

- Also possible to use a logical token ring structure
- When the algorithm is initialised process 0 has a token
- The process passes the ring to its neighbour in the token ring structure
- When the process receives the token it checks to see if it wants the resource
- If it does it utilises the resource and then releases the token
- A process will have to wait until every other process in the token utilises the resource in the worst case scenario

Assignment Project Exam Help

<https://tutorcs.com>

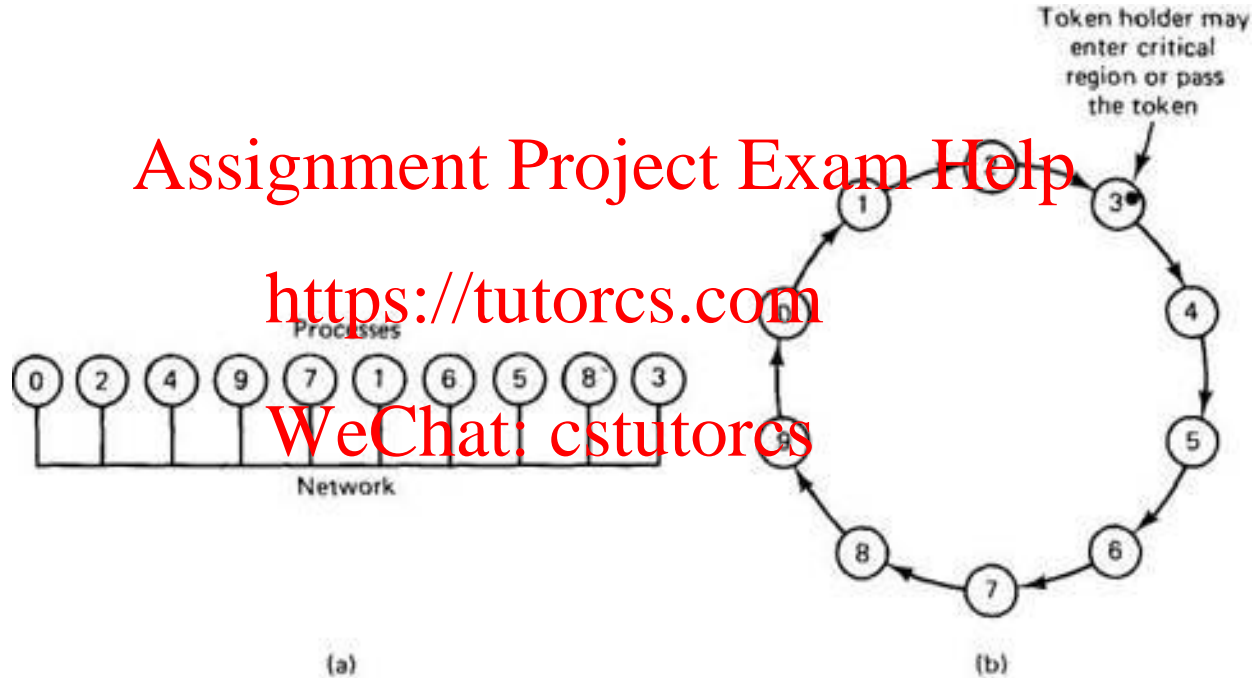
WeChat: cstutorcs

Token Ring Resource Allocation

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Distributed Operating Systems. Andrew Tanenbaum

Comparison of Resource Allocation Algorithm

Algorithm	Messages per Resource Allocation	Delay before Resource Allocation (in message times)	Potential Problems
Centralised	3	2	Coordinator Crash (Central Point of Failure)
Distributed	$2(n-1)$	$2(n-1)$	Crash of any process
Token Ring	$1 - \infty$	$0 - (n-1)$	Lost token, Process Crash

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Election Algorithms

- In the event of failure of a leader process in a centralized scheme and for other algorithms such as the Berkley Algorithm discussed earlier an election algorithm is necessary to select a new leader
- Election algorithms include:
 - Bully Algorithm
 - Ring Algorithm

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Bully Algorithm


- In the Bully algorithm each process is assigned a unique weight value
- If a process notices that the coordinator has stopped it sends election messages to processes which have a higher weight
- When a process receives an election message it replies to the election message. In addition if it has not sent messages to processes with a higher weight value it does so
- If the process receives an election message and it has sent messages it stops
- If after a certain time period a process has received no replies to its election message it determines that it has won the election and sends out a message to all processes indicating that it is the coordinator
- When a process recovers it will launch a new election

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Bully Algorithm



Distributed Operating Systems. Andrew Tanenbaum

Ring Algorithm


- In a ring algorithm when any process notices that the coordinator is down it sends an election message to the next process in the logical ring
- The election message includes its process number
- If the next process in the logical ring does not acknowledge the message then the it sends it the process after this in the logical ring until it receives and acknowledgement
- When a process receives an election message it adds its process number to the election message and forwards it to the next process in the ring
- When a process receives an election message it then sends a coordinator message around the ring which indicate who is the new coordinator (the highest process number) and the members of the logical ring
- When a process receives the coordinator message it has sent it removes the message as it has gone around the ring. This done to reduce overhead

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Ring Algorithm



Distributed Operating Systems. Andrew Tanenbaum